

OpenMP

This lesson deals with the OpenMP API which is popular for its handling of multiple threads.

We'll cover the following

- Private vs Shared variables
- Reduction
- Other OpenMP directives
- MPI
 - MPIReduce
 - An Example: Estimating pi using dartboard algorithm

[OpenMP](#) is an API that implements a multi-threaded, shared memory form of parallelism. It uses a set of compiler directives (statements that you add to your C code) that are incorporated at compile-time to generate a multi-threaded version of your code. You can think of Pthreads (above) as doing multi-threaded programming “by hand”, and OpenMP as a slightly more automated, higher-level API to make your program multithreaded. OpenMP takes care of many of the low-level details that you would normally have to implement yourself, if you were using Pthreads from the ground up.

Here is the general code structure of an OpenMP program:

```
#include <omp.h>

main () {

int var1, var2, var3;

Serial code
    .
    .
    .

Beginning of parallel section. Fork a team of threads.
Specify variable scoping

#pragma omp parallel private(var1, var2) shared(var3)
{
```



```
Parallel section executed by all threads
```

```
.  
.   
.
```

```
All threads join master thread and disband
```

```
}
```

```
Resume serial code
```

```
.  
.   
.
```

```
}
```

OpenMP

Private vs Shared variables

By using the `private()` and `shared()` directives, you can specify variables within the parallel region as being **shared**, i.e. visible and accessible by all threads simultaneously, or **private**, i.e. private to each thread, meaning each thread will have its own local copy. In the code example below for parallelizing a for loop, you can see that we specify the `thread_id` and `nloops` variables as `private`.

```
plg@wildebeest:~/Desktop$ gcc -fopenmp -o go go.c  
plg@wildebeest:~/Desktop$ ./go  
Thread 1 is adding its iterations (12500) to sum (0), total nloops is now 12500.  
Thread 4 is adding its iterations (12500) to sum (12500), total nloops is now 25000.  
Thread 0 is adding its iterations (12500) to sum (25000), total nloops is now 37500.  
Thread 5 is adding its iterations (12500) to sum (37500), total nloops is now 50000.  
Thread 3 is adding its iterations (12500) to sum (50000), total nloops is now 62500.  
Thread 6 is adding its iterations (12500) to sum (62500), total nloops is now 75000.  
Thread 2 is adding its iterations (12500) to sum (75000), total nloops is now 87500.  
Thread 7 is adding its iterations (12500) to sum (87500), total nloops is now 100000.  
Total # loop iterations is 100000
```

Reduction

Reduction refers to the process of combining the results of several sub-calculations into a final result. This is a very common paradigm (and indeed the so-called “map-reduce” framework used by Google and others is very popular). Indeed we used this paradigm in the code example above, where we used the “critical code” directive to accomplish this. The map-reduce paradigm is so common that OpenMP has a specific directive that allows you to more easily implement this.

```
#include <stdio.h>  
#include <omp.h>  
  
int main(int argc, char *argv[])
```

```

int main(int argc, char *argv[])
{
    int i, thread_id;
    int glob_nloops, priv_nloops;
    glob_nloops = 0;

    // parallelize this chunk of code
    #pragma omp parallel private(priv_nloops, thread_id) reduction(+:glob_nloops)
    {
        priv_nloops = 0;
        thread_id = omp_get_thread_num();

        // parallelize this for loop
        #pragma omp for
        for (i=0; i<100000; ++i)
        {
            ++priv_nloops;
        }
        glob_nloops += priv_nloops;
    }
    printf("Total # loop iterations is %d\n",
           glob_nloops);
    return 0;
}

```

```

plg@wildebeest:~/Desktop$ gcc -fopenmp -o go go.c
plg@wildebeest:~/Desktop$ ./go
Total # loop iterations is 100000

```



Other OpenMP directives

There are a host of other directives you can issue using OpenMP, see [here](#) for a list (wikipedia). Some other clauses of interest are:

- **barrier**: each thread will wait until all threads have reached this point in the code, before proceeding
- **nowait**: threads will not wait until everybody is finished
- **schedule(type, chunk)** allows you to specify how tasks are spawned out to threads in a for loop. There are three types of scheduling you can specify
- **if**: allows you to parallelize only if a certain condition is met
- ... and a host of others

MPI

The [Message Passing Interface](#) (MPI) is a standard defining core syntax and semantics of library routines that can be used to implement parallel programming in C (and in other languages as well). There are several implementations of MPI such as [Open MPI](#), [MPICH2](#) and [LAM/MPI](#).

In the context of this tutorial, we will think of MPI in terms of its role in parallelizing

In the context of this tutorial, you can think of MPI, in terms of its complexity, scope and control, as sitting in between programming with Pthreads, and using a high-level API such as OpenMP.

The MPI interface allows you to manage allocation, communication, and synchronization of a set of processes that are mapped onto multiple nodes, where each node can be a core within a single CPU, or CPUs within a single machine, or even across multiple machines (as long as they are networked together).

One context where MPI shines in particular is the ability to easily take advantage not just of multiple cores on a single machine, but to run programs on clusters of several machines. Even if you don't have a dedicated cluster, you could still write a program using MPI that could run your program in parallel, across any collection of computers, as long as they are networked together. Just make sure to ask permission before you load up your lab-mate's computer's CPU(s) with your computational tasks!

Here is a basic MPI program that simply writes a message to the screen indicating which node is running.

```
// mpicc go_mpi.c -o go_mpi
// mpirun -n 4 go_mpi

#include <stdio.h>
#include <mpi.h>

int main(int argc, char *argv[])
{
    int myrank, nprocs;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);

    printf("I am node %d of %d\n", myrank, nprocs);

    MPI_Finalize();
    return 0;
}
```

```
plg@wildebeest:~/Desktop$ mpicc go_mpi.c -o go_mpi
plg@wildebeest:~/Desktop$ mpirun -n 4 go_mpi
I am node 0 of 4
I am node 2 of 4
I am node 1 of 4
I am node 3 of 4
```

The basic design pattern of an MPI program is that the **same code** is sent to all

nodes for execution. It's by using the `MPI_Comm_rank()` function that you can

determine which node is running, and (if needed) act differently. The `MPI_Comm_size()` function will tell you how many nodes there are in total.

MPI programs need to be compiled using `mpicc`, and need to be run using `mpirun` with a flag indicating the number of processors to spawn (4, in the above example).

MPI Reduce

We saw with OpenMP that we can use a **reduce** directive to sum values across all threads. A similar function exists in MPI called `MPI_Reduce()`.

An Example: Estimating pi using dartboard algorithm

```
// Estimating pi using the dartboard algorithm
// All processes contribute to the calculation, with the
// master process averaging the values for pi.
// We then use mpc_reduce to collect the results
//
// mpicc -o go mpi_pi_reduce.c
// mpirun -n 8 go

#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>

#define MASTER 0          // task ID of master task
#define NDARTS 1000       // # dart throws per round
#define NROUNDS 10        // # of rounds of dart throwing

// our function for throwing darts and estimating pi
double dartboard(int ndarts)
{
    double x, y, r, pi;
    int n, hits;
    hits = 0;

    // throw darts
    for (n = 1; n <= ndarts; n++) {
        // (x,y) are random between -1 and 1
        r = (double)random()/RAND_MAX;
        x = (2.0 * r) - 1.0;
        r = (double)random()/RAND_MAX;
        y = (2.0 * r) - 1.0;
        // if our random dart landed inside the unit circle, increment the score
        if (((x*x) + (y*y)) <= 1.0) {
            hits++;
        }
    }

    // estimate pi
    pi = 4.0 * (double)hits / (double)ndarts;
```

```

return(pi);
}

// the main program
int main (int argc, char *argv[])
{
    double my_pi, pi_sum, pi_est, mean_pi, err;
    int task_id, n_tasks, rc, i;
    MPI_Status status;

    // Obtain number of tasks and task ID
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&n_tasks);
    MPI_Comm_rank(MPI_COMM_WORLD,&task_id);
    // printf ("task %d of %d reporting for duty...\n", task_id, n_tasks);

    // different seed for random number generator for each task
    srandom (task_id);

    mean_pi = 0.0;
    for (i=0; i<NROUNDS; i++) {
        // all tasks will execute dartboard() to calculate their own estimate of pi
        my_pi = dartboard(NDARTS);

        // now we use MPI_Reduce() to sum values of my_pi across all tasks
        // the master process (id=MASTER) will store the accumulated value
        // in pi_sum. We tell MPI_Reduce() to sum by passing it
        // the MPI_SUM value (define in mpi.h)
        rc = MPI_Reduce(&my_pi, &pi_sum, 1, MPI_DOUBLE, MPI_SUM,
                        MASTER, MPI_COMM_WORLD);

        // now, IF WE ARE THE MASTER process, we will compute the mean
        if (task_id == MASTER) {
            pi_est = pi_sum / n_tasks;
            mean_pi = ( (mean_pi * i) + pi_est ) / (i + 1); // running average
            err = mean_pi - 3.14159265358979323846;
            printf("%d throws: mean_pi %.12f: error %.12f\n",
                (NDARTS * (i + 1)), mean_pi, err);
        }
    }
    if (task_id == MASTER)
        printf ("PS, the real value of pi is about 3.14159265358979323846\n");

    MPI_Finalize();
    return 0;
}

```

Here we run it with just one parallel process:

```

plg@wildebeest:~/Desktop/mpi$ time mpirun -n 1 go
1000 throws: mean_pi 3.088000000000: error -0.053592653590
2000 throws: mean_pi 3.104000000000: error -0.037592653590
3000 throws: mean_pi 3.101333333333: error -0.040259320256
4000 throws: mean_pi 3.120000000000: error -0.021592653590
5000 throws: mean_pi 3.124800000000: error -0.016792653590
6000 throws: mean_pi 3.127333333333: error -0.014259320256
7000 throws: mean_pi 3.134285714286: error -0.007306939304
8000 throws: mean_pi 3.128500000000: error -0.013092653590
9000 throws: mean_pi 3.132444444444: error -0.009148209145

```



```
10000 throws: mean_pi 3.119600000000: error -0.021992653590  
PS, the real value of pi is about 3.14159265358979323846
```

```
real    0m0.032s  
user    0m0.020s  
sys 0m0.012s
```

Now let's run it with 4:

```
plg@wildebeest:~/Desktop/mpi$ time mpirun -n 4 go  
1000 throws: mean_pi 3.105000000000: error -0.036592653590  
2000 throws: mean_pi 3.122500000000: error -0.019092653590  
3000 throws: mean_pi 3.122000000000: error -0.019592653590  
4000 throws: mean_pi 3.137750000000: error -0.003842653590  
5000 throws: mean_pi 3.143600000000: error 0.002007346410  
6000 throws: mean_pi 3.140166666667: error -0.001425986923  
7000 throws: mean_pi 3.142000000000: error 0.000407346410  
8000 throws: mean_pi 3.140250000000: error -0.001342653590  
9000 throws: mean_pi 3.136666666667: error -0.004925986923  
10000 throws: mean_pi 3.135000000000: error -0.006592653590  
PS, the real value of pi is about 3.14159265358979323846
```

```
real    0m0.034s  
user    0m0.044s  
sys 0m0.024s
```

We see the final error is much reduced. Each of the 4 processes (which are parallelized across the cores of my CPU) contributes an estimate of pi, which are then averaged by the master process to come up with the final estimate of pi.

After OpenMP, we'll move a layer above to define what tools such as this are. Such APIs follow the Message Passing Interface or MPI. The next lesson deals with this concept in detail.