

# Verifying Jenkins X Boot Installation

In this lesson, we will verify the installation done by Jenkins X Boot by checking pipelines, Namespaces, activities and creating a new quickstart project.

## We'll cover the following

- Checking the pipelines
- Triggering the dev pipeline
- Checking the Namespaces present in the cluster
- Checking the environments
- Creating a new quickstart project
- Inspecting the activities and Namespaces

## Checking the pipelines #

Let's take a quick look at the pipelines currently active in our cluster.

```
jx get pipelines --output yaml
```

The output is as follows.

```
- vfarcic/environment-jx-boot-dev/master  
- vfarcic/environment-jx-boot-production/master  
- vfarcic/environment-jx-boot-staging/master
```

We are already used to working with `production` and `staging` pipelines. What is new is the `dev` pipeline. That is the one we just executed locally. It is now available in the cluster as well, and we should be able to trigger it by pushing a change to the associated repository. Let's test that.

## Triggering the `dev` pipeline #

We'll explore the Jenkins X upgrade process later. For now, we just want to see whether the `dev` repository is indeed triggering pipeline activities. We'll do that by making a trivial change to the `README.md` file.

```
echo "A trivial change" \
  | tee -a README.md

git add .

git commit -m "A trivial change"

git push

jx get activities \
  --filter environment-$CLUSTER_NAME-dev \
  --watch
```

We pushed the changes to GitHub and started watching the activities of the `dev` pipeline. The output, when the activity is finished, should be as follows.

STEP	STARTED AGO	DURATION	STATUS
vfarcic/environment-jx-boot-dev/master #1	4m10s	3m52s	Succeeded
release	4m10s	3m52s	Succeeded
Credential Initializer 7jh9t	4m10s	0s	Succeeded
Working Dir Initializer Tr2wz	4m10s	2s	Succeeded
Place Tools	4m8s	2s	Succeeded
Git Source Vfarcic Environment Jx...	4m6s	36s	Succeeded <a href="https://github.com/vfarcic">https://github.com/vfarcic</a>
Git Merge	3m30s	1s	Succeeded
Validate Git	3m29s	1s	Succeeded
Verify Preinstall	3m28s	26s	Succeeded
Install Jx Crds	3m2s	10s	Succeeded
Install Velero	2m52s	12s	Succeeded
Install Velero Backups	2m40s	2s	Succeeded
Install Nginx Controller	2m38s	16s	Succeeded
Create Install Values	2m22s	0s	Succeeded
Install External Dns	2m22s	16s	Succeeded
Install Cert Manager Crds	2m6s	0s	Succeeded
Install Cert Manager	2m6s	16s	Succeeded
Install Acme Issuer And Certificate	1m50s	2s	Succeeded
Install Vault	1m48s	8s	Succeeded
Create Helm Values	1m40s	3s	Succeeded
Install Jenkins X	1m37s	1m3s	Succeeded
Verify Jenkins X Environment	34s	5s	Succeeded
Install Repositories	29s	5s	Succeeded
Install Pipelines	24s	1s	Succeeded
Update Webhooks	23s	4s	Succeeded
Verify Installation	19s	1s	Succeeded

That's the first activity (`#1`) of the `dev` pipeline. To be more precise, it is the second one (the first was executed locally) but, from the perspective of Jenkins X inside the cluster, which did not exist at the time, that is the first activity. Those are the steps of Jenkins X Boot running inside our cluster.

We won't go through the changes that were created by that activity since there are none. We did not modify any of the files that matter. We already used those same files when we executed Jenkins X Boot locally. All we did was push the change of the `README` file from the local repository to GitHub. That triggered a webhook that

the `README` file from the local repository to GitHub. That triggered a webhook that notified the cluster that there are some changes to the remote repo. As a result, the first in-cluster activity was executed.

The reason I showed you that activity was not due to an expectation of seeing some change applied to the cluster (there were none), but to demonstrate that, from now on, we should let Jenkins X running inside our Kubernetes cluster handle changes to the `dev` repository, instead of running `jx boot` locally. That will come in handy later on when we explore how to upgrade or change our Jenkins X setup.

Please press `ctrl+c` to stop watching the activity.

## Checking the Namespaces present in the cluster #

Let's take a look at the Namespaces we have in our cluster.

```
kubectl get namespaces
```

The output is as follows.

NAME	STATUS	AGE
cert-manager	Active	28m
default	Active	74m
jx	Active	34m
kube-public	Active	74m
kube-system	Active	74m
velero	Active	33m

As you can see, there are no Namespaces for staging and production environments. **Does that mean that we do not get them with Jenkins X Boot?**

Unlike other types of setup, the Boot creates environments lazily. That means that they are not created in advance, but rather when used for the first time. In other words, the `jx-staging` Namespace will be created the first time we deploy something to the staging environment. The same logic is applied to any other environment, the production included.

## Checking the environments #

To put your mind at ease, we can output the environments and confirm that staging and production were indeed created.

The output is as follows.

NAME	LABEL	KIND	PROMOTE	NAMESPACE	ORDER	CLUSTER	SOURCE
dev	Development	Development	Never	jx	0		<a href="https://github.com/vfarcic/">https://github.com/vfarcic/</a>
staging	Staging	Permanent	Auto	jx-staging	100		<a href="https://github.com/vfarcic/">https://github.com/vfarcic/</a>
production	Production	Permanent	Manual	jx-production	200		<a href="https://github.com/vfarcic/">https://github.com/vfarcic/</a>

`Staging` and `production` environments do indeed exist, and they are associated with corresponding Git repositories, even though Kubernetes Namespaces were not yet created.

While we're on the subject of environments, we can see something that did not exist in the previous setups. What's new, in this output, is that the `dev` environment also has a repo set as the source. That was not the case when we were creating clusters with `jx create cluster` or installing Jenkins X with `jx install`. In the past, the `dev` environment was not managed by GitOps principles. Now it is.

The associated `source` is the repository we forked, and it contains the specification of the whole `dev` environment and a few other things. From now on, if we need to modify Jenkins X or any other component of the platform, all we have to do is push a change to the associated repository. A pipeline will take care of converging the actual with the desired state. From this moment onward, everything except infrastructure is managed through GitOps. I intentionally said “*except infrastructure*” because I did not show you explicitly how to create a pipeline that will execute **Terraform** or whichever tool you're using to manage your infra. However, that is not an excuse for you not to do it.

By now, you should have enough knowledge to automate that last piece of the puzzle by creating a Git Repository with `jenkins-x.yml` and importing it into Jenkins X. If you do that, everything will be using GitOps, and everything will be automated, except writing code and pushing it to Git. Do not take my unwillingness to force a specific tool (e.g., **Terraform**) as a sign that you shouldn't walk that last mile.

## Creating a new quickstart project #

To be on the safe side, we'll create a new quickstart project and confirm that the system is behaving correctly. We still need to verify that lazy creation of

system is behaving correctly. We still need to verify that the creation of environments works as expected.

```
cd ..  
  
jx create quickstart \  
  --filter golang-http
```

We went back from the local copy of the `dev` repository and started the process of creating a new `quickstart`.

You might be wondering why we didn't use the `--batch-mode` as we did countless times before. This time, we installed Jenkins X using a different method, and the local cache does not have the information it needs to create a `quickstart` automatically. So, we'll need to answer a series of questions, just as we did at the beginning of the book. That is only a temporary issue. The moment we create the first quickstart manually, the information will be stored locally, and every consecutive attempt to create new quickstarts can be done with the `--batch-mode`, as long as we're using the same laptop.

Please answer the questions any way you like. Just bear in mind that I expect you to name the repository `jx-boot`. If you call it differently, you'll have to modify the commands that follow.

## Inspecting the activities and Namespaces #

All that's left now is to wait until the activity of the pipeline created by the quickstart process is finished.

```
jx get activity \  
  --filter jx-boot/master \  
  --watch
```

Please press `ctrl+c` to stop watching the activity once it's finished.

We'll also confirm that the activity of the staging environment is finished as well.

```
jx get activity \  
  --filter environment-$CLUSTER_NAME-staging/master \  
  --watch
```

You know what to do. Wait until the new activity is finished and press `ctrl+c` to stop watching.

Let's retrieve the Namespaces again.

```
kubectl get namespaces
```

The output is as follows.

NAME	STATUS	AGE
cert-manager	Active	34m
default	Active	80m
jx	Active	40m
jx-staging	Active	55s
kube-public	Active	80m
kube-system	Active	80m
velero	Active	39m

As you can see, the `jx-staging` Namespace was created the first time an application was deployed to the associated environment, not when we installed the platform. The `jx-production` Namespace is still missing since we did not yet promote anything to production. Once we do, that Namespace will be created as well.

I won't bother you with the commands that would confirm that the application is accessible from outside the cluster, that we can create preview (pull request) environments, and that we can promote releases to production. I'm sure that, by now, you know how to do all that by yourself.

---

The next lesson will wrap up this chapter and give instructions to free up the used resources.